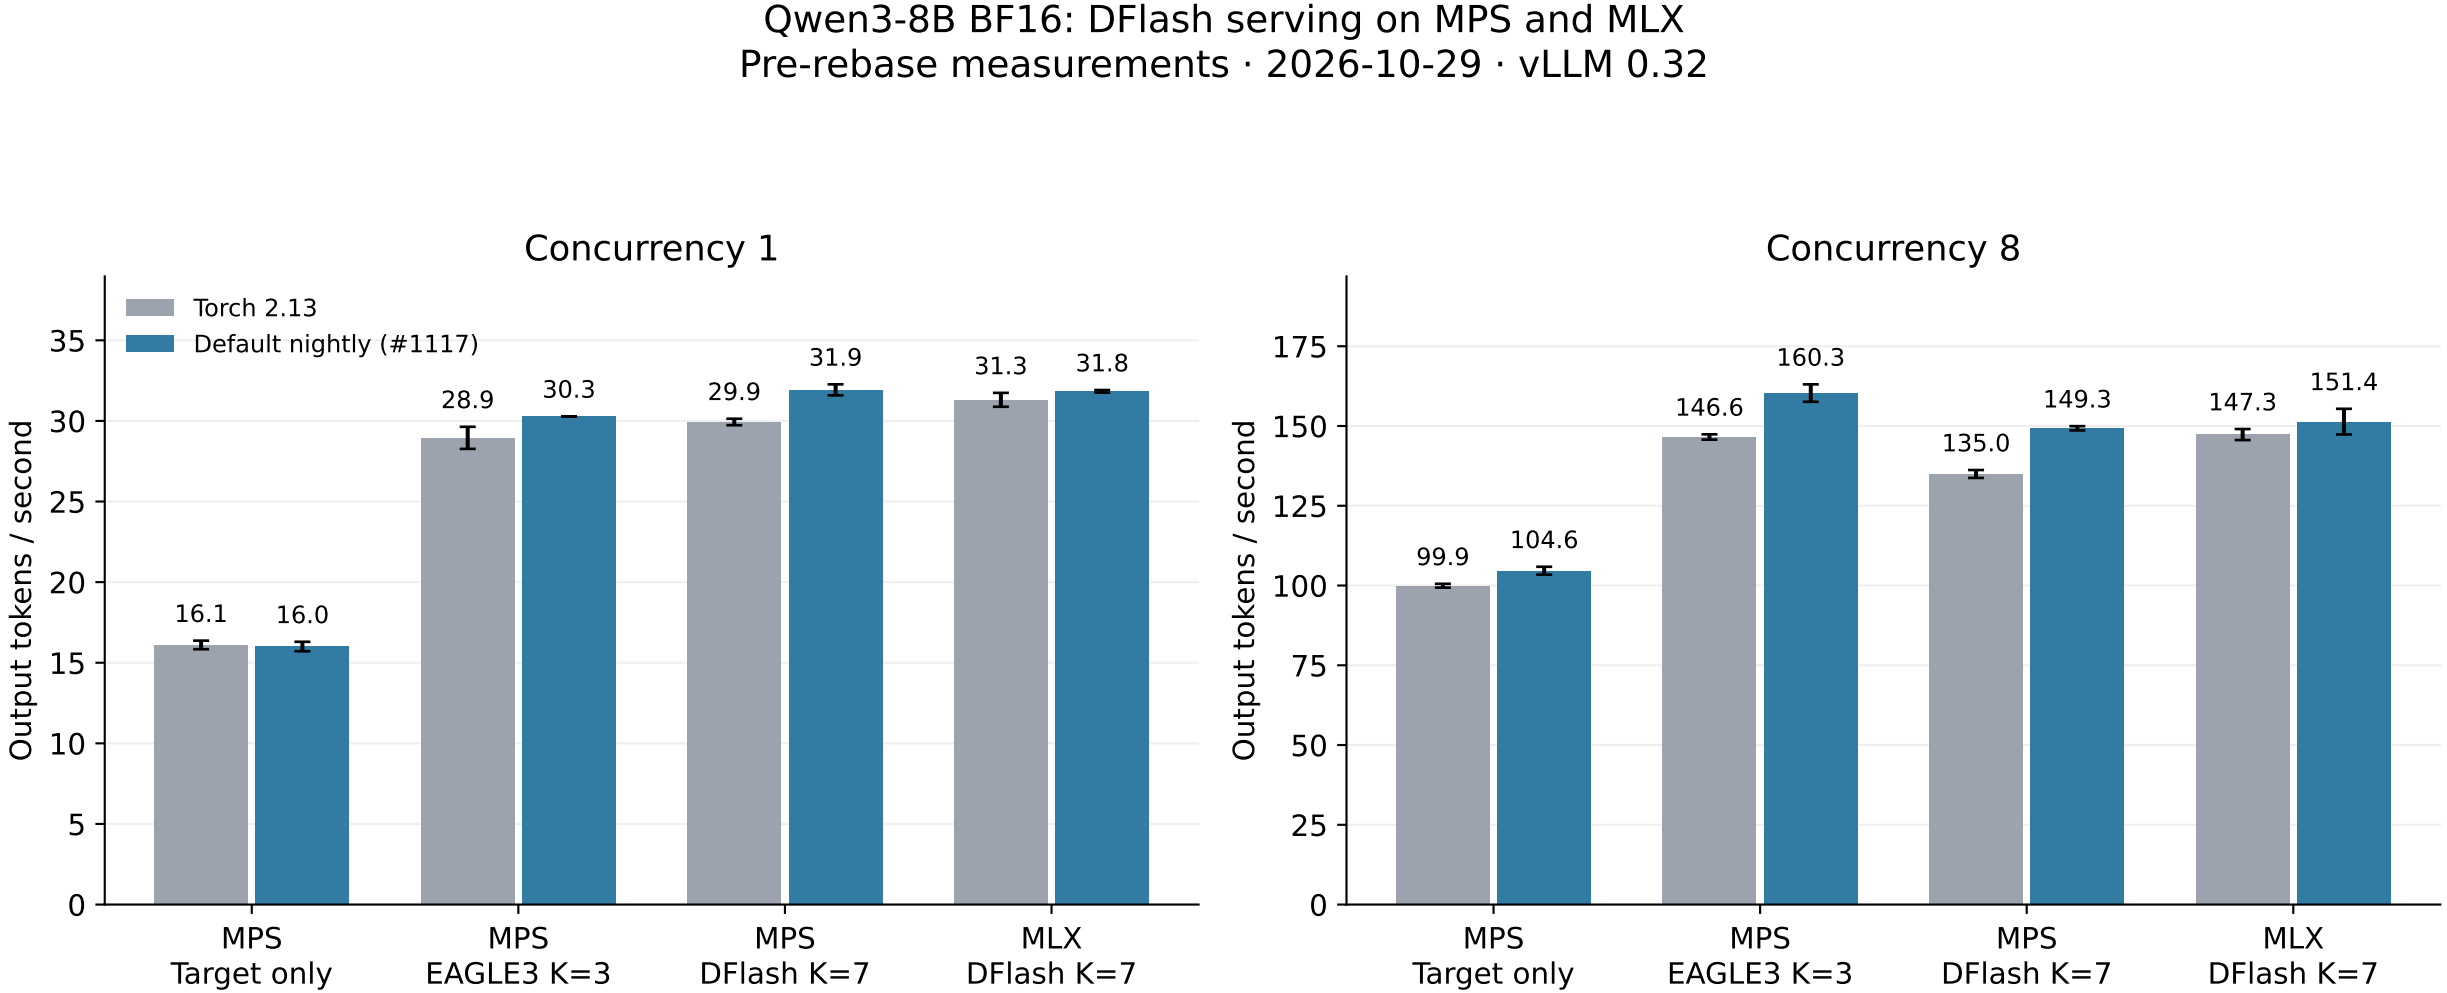
Qwen3-8B BF16: DFlash serving on MPS and MLX
Pre-rebase measurements · 2026-10-29 · vLLM 0.32
Concurrency 1
Concurrency 8
Torch 2.13
Default nightly (#1117)
Output tokens / second
16.1
16.0
28.9
30.3
29.9
31.9
31.3
31.8
99.9
104.6
146.6
160.3
135.0
149.3
147.3
151.4
MPS Target only
MPS EAGLE3 K=3
MPS DFlash K=7
MLX DFlash K=7
36 fixed Speculator Bench prompts × 128 output tokens · Two runs per setting; reversed configuration order
Bars: mean; whiskers: range · Same z-lab DFlash draft · Nightly: 2.16.0.dev20261013 · M5 Pro, 20 GPU cores, 64 GiB